

Complete Terminal of the Running Agent

Last login: Tue Jun 2 18:13:56 on ttys000

asifkhan@ASIFs-MacBook-Air ~ % ollama list

NAME	ID	SIZE	MODIFIED
------	----	------	----------

mistral:latest	6577803aa9a0	4.4 GB	18 hours ago
----------------	--------------	--------	--------------

asifkhan@ASIFs-MacBook-Air ~ % ollama pull mistral:latest

pulling manifest

pulling f5074b1221da: 100%

4.4 GB

pulling 43070e2d4e53: 100%

11 KB

pulling 1ff5b64b61b9: 100%

799 B

pulling ed11eda7790d: 100%

30 B

pulling 1064e17101bd: 100%

487 B

verifying sha256 digest

writing manifest

success

asifkhan@ASIFs-MacBook-Air ~ % cd ~/Desktop/AI_ML/Agentic_AI_Deep_Research

asifkhan@ASIFs-MacBook-Air Agentic_AI_Deep_Research % source myenv/bin/activate

(myenv) asifkhan@ASIFs-MacBook-Air Agentic_AI_Deep_Research % wc -l predictions/
*.jsonl

30 predictions/baseline.jsonl

0 predictions/full_agent.jsonl

30 predictions/no_hybrid.jsonl

30 predictions/no_planner.jsonl

30 predictions/no_reflector.jsonl

1 predictions/no_verifier.jsonl

121 total

```
(myenv) asifkhan@ASIFs-MacBook-Air Agentic_AI_Deep_Research % python main.py \
--config full_agent \
--questions eval/questions.jsonl \
--output predictions/full_agent.jsonl
```

Connecting to arXiv API via url: [http://export.arxiv.org/api/query?](http://export.arxiv.org/api/query?search_query=cat:cs.CL+OR+cat:cs.AI+OR+cat:cs.LG&start=0&max_results=500&sortBy=submittedDate&sortOrder=descending)

[search_query=cat:cs.CL+OR+cat:cs.AI+OR+cat:cs.LG&start=0&max_results=500&sortBy=submittedDate&sortOrder=descending](http://export.arxiv.org/api/query?search_query=cat:cs.CL+OR+cat:cs.AI+OR+cat:cs.LG&start=0&max_results=500&sortBy=submittedDate&sortOrder=descending)

-> Processing paper match [2606.02578]: Mitigating Perceptual Judgment Bias in Multimodal ...

-> Processing paper match [2606.02576]: ProtoAda: Prototype-Guided Adaptive Adapter Expans...

-> Processing paper match [2606.02569]: AdaCodec: A Predictive Visual Code for Video MLLMs...

-> Processing paper match [2606.02568]: ClinEnv: An Interactive Multi-Stage Long Horizon E...

-> Processing paper match [2606.02563]: IntraShuffler: A Privacy Preserving Framework for ...

-> Processing paper match [2606.02562]: Permissive Safety Through Trusted Inference: Verif...

-> Processing paper match [2606.02559]: From Layers to Submodules: Rethinking Granularity ...

-> Processing paper match [2606.02556]: HERO'S JOURNEY: Testing Complex Rule Induction wit...

-> Processing paper match [2606.02552]: Modeling Depth Ambiguity: A Mixture-Density Repres...

-> Processing paper match [2606.02548]: SN-WER: Script-Normalized WER for Multi-Script Ind...

-> Processing paper match [2606.02545]: Transferable Self-Harm Surveillance from Emergency...

-> Processing paper match [2606.02544]: SimSD: Simple Speculative Decoding in Diffusion La...

-> Processing paper match [2606.02540]: SkillHarm: Lifecycle-Aware Skill-Based Attacks via...

-> Processing paper match [2606.02536]: Tracking the Behavioral Trajectories of Adapting A...

-> Processing paper match [2606.02530]: SafeSteer: Localized On-Policy Distillation for Ef...

-> Processing paper match [2606.02528]: Auditing Asset-Specific Preferences in Financial L...

-> Processing paper match [2606.02526]: Why Not Hyperparameter-Friendly Optimisation? A Mo...

-> Processing paper match [2606.02523]: FigSIM: A Dataset for Fine-grained Suicide Severit...

-> Processing paper match [2606.02522]: Moment-Video: Diagnosing Temporal Fidelity of Vide...

-> Processing paper match [2606.02521]: Drifting Preference Optimization for One-Step Gene...

-> Processing paper match [2606.02515]: A Biconvex Formulation for Stable Transport of Mix...

-> Processing paper match [2606.02509]: When Rating Scales Fall Short: LLM-Assisted Discov...

-> Processing paper match [2606.02507]: Towards Automated Discovery: A Review of Generativ...

-> Processing paper match [2606.02502]: CRAM: Centroid-Routing and Adaptive MoE for Multim...

-> Processing paper match [2606.02497]: Bridging the Last Mile of Time Series Forecasting ...

-> Processing paper match [2606.02494]: Monitoring Agentic Systems Before They're Reliable...

-> Processing paper match [2606.02493]: Not What, But How: A Communicative Audit of LLM Re...

-> Processing paper match [2606.02490]: Expressivity of congruence-based architectures for...

-> Processing paper match [2606.02488]: RASER: Recoverability-Aware Selective Escalation R...

-> Processing paper match [2606.02487]: Towards Multidisciplinary Summarization of Hospita...

-> Processing paper match [2606.02484]: Iteris: Agentic Research Loops for Computational M...

-> Processing paper match [2606.02483]: Ghost Tool Calls: Issue-Time Privacy for Speculati...

-> Processing paper match [2606.02475]: Physics-Informed Residuals for Adaptive Mesh Refin...

-> Processing paper match [2606.02470]: MCP-Persona: Benchmarking LLM Agents on Real-World...

-> Processing paper match [2606.02465]: Learning When to Translate for Multilingual Reason...

-> Processing paper match [2606.02463]: MASER: Modality-Adaptive Specialist Routing for Em...

-> Processing paper match [2606.02461]: AGENTCL: Toward Rigorous Evaluation of Continual L...

-> Processing paper match [2606.02458]: Beyond One-shot: AI Agents for Learning in Field E...

-> Processing paper match [2606.02455]: Speculative Sampling For Faster Molecular Dynamics...

-> Processing paper match [2606.02453]: Initialization is Half the Battle: Generating Dive...

-> Processing paper match [2606.02449]: HLL: Can Agents Cross Humanity's Last Line of Veri...

-> Processing paper match [2606.02444]: Food Noise & False Safety: A Systematic Evaluation...

-> Processing paper match [2606.02443]: PaSBench-Video: A Streaming Video Benchmark for Pr...

-> Processing paper match [2606.02438]: LLM-Evolved Pattern Generators for Optimal Classic...

-> Processing paper match [2606.02437]: On the Scaling of PEFT: Towards Million Personal M...

-> Processing paper match [2606.02433]: ODTQA-FoRe: An Open-Domain Tabular Question Answer...

-> Processing paper match [2606.02434]: Bridging the Sim-to-Real Gap in Semiconductor Visu...

-> Processing paper match [2606.02430]: Not All Errors Are Equal: A Systematic Study of Er...

-> Processing paper match [2606.02427]: Spectral Audit of In-Context Operator Networks...

-> Processing paper match [2606.02424]: GC-MoE: Genomics-Guided Cell-Type-Specific Mixture...

-> Processing paper match [2606.02423]: Investigating and Alleviating Harm Amplification i...

-> Processing paper match [2606.02418]: Evolutionary Discovery of Bivariate Bicycle Codes ...

-> Processing paper match [2606.02404]: K-BrowseComp: A Web Browsing Agent Benchmark Groun...

-> Processing paper match [2606.02403]: AutoForest: Automatically Generating Forest Plots ...

-> Processing paper match [2606.02398]: A Local Perturbation Theory for Cross-Domain Inter...

-> Processing paper match [2606.02388]: Policy and World Modeling Co-Training for Language...

-> Processing paper match [2606.02386]: AgentPLM: Agentic Protein Language Models with Rea...

-> Processing paper match [2606.02385]: How Optimality Structures Sparse Dictionaries: A T...

-> Processing paper match [2606.02384]: TabPrep: Closing the Feature Engineering Gap in Ta...

-> Processing paper match [2606.02381]: A Mathematical Conflict Framework for Contextual D...

-> Processing paper match [2606.02380]: SPADE-Bench: Evaluating Spontaneous Strategic Dece...

-> Processing paper match [2606.02378]: When Do Attention Circuits Form? Developmental Tra...

-> Processing paper match [2606.02375]: WAXAL-NET: Finetuned Edge ASR Across 19 African La...

-> Processing paper match [2606.02374]: Spatial Representation Learning Beyond Pixels: Uni...

-> Processing paper match [2606.02373]: Harness-1: Reinforcement Learning for Search Agent...

-> Processing paper match [2606.02372]: COMAP: Co-Evolving World Models and Agent Policies...

-> Processing paper match [2606.02365]: FOAM: Frequency and Operator Error-Based Adaptive ...

-> Processing paper match [2606.02363]: Minimax-Optimal Policy Regret in Partially Observa...

-> Processing paper match [2606.02359]: MOC: Multi-Order Communication in LLM-based Multi-...

-> Processing paper match [2606.02357]: Do Multimodal Agents Really Benefit from Tool Use?...

-> Processing paper match [2606.02355]: SIRI: Self-Internalizing Reinforcement Learning wi...

-> Processing paper match [2606.02351]: Local Preferential Bayesian Optimization...

-> Processing paper match [2606.02345]: Doing well with less! On Sampling Techniques for E...

-> Processing paper match [2606.02341]: Parameter-efficient Dual-encoder Architecture with...

-> Processing paper match [2606.02339]: Entropy Minimization without Model Collapse: Mitig...

-> Processing paper match [2606.02337]: Coordination Graphs for Constrained Multi-Agent Re...

-> Processing paper match [2606.02332]: Forget Attention: Importance-Aware Attention Is Al...

-> Processing paper match [2606.02331]: Hallucination-Aware Diffusion Sampling for Inverse...

-> Processing paper match [2606.02328]: Riemannian Gradient Descent for Low-Rank Architect...

-> Processing paper match [2606.02326]: Repair Before Veto: Repair-Augmented Constraint Le...

-> Processing paper match [2606.02322]: Repurposing Adversarial Perturbations for Continua...

-> Processing paper match [2606.02320]: TVIR: Building Deep Research Agents Towards Text--...

-> Processing paper match [2606.02310]: Deep Learning for Remote Sensing to Improve Flood ...

-> Processing paper match [2606.02309]: Measurement Geometry and Design for Trustworthy Ge...

-> Processing paper match [2606.02304]: Unified Context Evolution for LLM Agents...

-> Processing paper match [2606.02302]: SeClaw: Spec-Driven Security Task Synthesis for Ev...

-> Processing paper match [2606.02301]: Quantitative Movement Testing: Measuring Patient M...

-> Processing paper match [2606.02300]: Beyond Isolated Behaviors: Hierarchical User Model...

-> Processing paper match [2606.02294]: Regularized Large Neighborhood Search...

-> Processing paper match [2606.02293]: AI as a Tool for Simulation-Based Experiments in L...

-> Processing paper match [2606.02289]: DECK: A Consistency x Confidence Taxonomy of LLM H...

-> Processing paper match [2606.02288]: Massive Spikes in LLMs are Bias Vectors: Mechanist...

-> Processing paper match [2606.02287]: CityTrajBench: A Unified Benchmark for City-Scale ...

-> Processing paper match [2606.02282]: POIROT: Interrogating Agents for Failure Detection...

-> Processing paper match [2606.02278]: Physics-Guided Recurrent State-Space Neural Networ...

-> Processing paper match [2606.02276]: Cross-modal linkage risk in clinical vision-langua...

-> Processing paper match [2606.02267]: A combination of noise and bilateral filters achie...

-> Processing paper match [2606.02256]: ArrythML: An Autoencoder-Based TinyML Approach for...

-> Processing paper match [2606.02255]: Who Annotates in NLP? A Large-scale Assessment of ...

-> Processing paper match [2606.02252]: ResMerge: Residual-based Spectral Merging of Large...

-> Processing paper match [2606.02253]: CEON: Circular Economy Ontology Network...

-> Processing paper match [2606.02251]: FW-NKF: Frequency-Weighted Neural Kalman Filters...

-> Processing paper match [2606.02248]: Geometric Latent Reasoning Induces Shorter Generat...

-> Processing paper match [2606.02247]: ShapLEIG: Bayesian Experimental Design for Shapley...

-> Processing paper match [2606.02245]: When Knowledge Is Not Free: Cost-Aware Evidence Se...

-> Processing paper match [2606.02242]: Towards Resolving Optimization Conflicts Between I...

-> Processing paper match [2606.02241]: BlockGen: Flexible Blockwise Sequence Modeling wit...

-> Processing paper match [2606.02240]: AgentRedBench: Dynamic Redteaming and Integration-...

-> Processing paper match [2606.02237]: Why Are DMD Students Lazy? Understanding the Copyi...

-> Processing paper match [2606.02232]: A Doeblin-Anchored Contrastive Chart for Learning ...

-> Processing paper match [2606.02231]: Identifiable Markov Switching Models with Instanta...

-> Processing paper match [2606.02228]: Bayesian meta-learning for modeling Alzheimer's di...

-> Processing paper match [2606.02223]: Network Learning with Semi-relaxed Gromov-Wasserst...

-> Processing paper match [2606.02221]: CORE-MTL: Rethinking Gradient Balancing via Causal...

-> Processing paper match [2606.02218]: Faster Synchronous On-Policy RL via Straggler-Awar...

-> Processing paper match [2606.02215]: Better with Experience: Self-Evolving LLM Agents f...

-> Processing paper match [2606.02214]: Do Gender Cues Affect LLM Value Trade-offs? Eviden...

-> Processing paper match [2606.02211]: Consistency Training while Mitigating Obfuscation ...

-> Processing paper match [2606.02204]: Cross-Environment Neural Reranking for Sample-Effi...

-> Processing paper match [2606.02198]: Model Multiplicity and Predictive Arbitrariness in...

-> Processing paper match [2606.02194]: Coherent Off-Policy Improvement of Large Behavior ...

-> Processing paper match [2606.02184]: The Ghost Couple: Correlated LLM Name Priors and T...

-> Processing paper match [2606.02179]: On the Generalization in Topology Optimization via...

-> Processing paper match [2606.02178]: Order within Chaos: Capturing Intrinsic Energy Ano...

-> Processing paper match [2606.02177]: Low-Pass Flow Matching...

-> Processing paper match [2606.02172]: Closing the Alignment-Maturity Gap in Federated Pr...

-> Processing paper match [2606.02170]: CRAFTQA: A Code-Driven Adaptive Framework for Comp...

-> Processing paper match [2606.02168]: Disentanglement-Based Equivariant Learning for Com...

-> Processing paper match [2606.02167]: From Capability Models to Automated Planning: An A...

-> Processing paper match [2606.02166]: EEG-FuseFormer: A Transformer-Driven Feature Fusio...

-> Processing paper match [2606.02163]: An Abstract Worlds Semantic Framework for Belief C...

-> Processing paper match [2606.02162]: Multimodal Approaches for Visually-Rich Document T...

-> Processing paper match [2606.02161]: InfoMerge: Information-aware Token Compression for...

-> Processing paper match [2606.02158]: On the Saliency of Low-Probability Tokens for AI-G...

-> Processing paper match [2606.02156]: Predicting the risk of colorectal anastomotic leak...

-> Processing paper match [2606.02151]: S3TS: Stochastic Scenario-Structured Tree Search f...

-> Processing paper match [2606.02147]: Multilingual Idioms in Sentences and Conversations...

-> Processing paper match [2606.02145]: Hybrid Neural Ordinary Differential Equations for ...

-> Processing paper match [2606.02142]: TimeBlocks: Foundational and Continual Time-Series...

-> Processing paper match [2606.02138]: VLBM: Variational Latent Basis Modeling for OOD Ro...

-> Processing paper match [2606.02136]: Edge-aware Decoding for Neural Asymmetric Routing...

-> Processing paper match [2606.02134]: Rethinking Evaluation Paradigms in IBP-based Certi...

-> Processing paper match [2606.02133]: Variational Learning for Insertion-based Generatio...

-> Processing paper match [2606.02132]: Learning When Not to Act: Mitigating Tool Abuse in...

-> Processing paper match [2606.02120]: Understanding-Enhanced Model Collaboration for Lon...

-> Processing paper match [2606.02119]: How Hard Can It Be? Hardness-Aware Multi-Objective...

-> Processing paper match [2606.02117]: ProbRes: Volatility Learning for Probabilistic Tim...

-> Processing paper match [2606.02115]: Error Bounds for a Diffusion Model-Based Drift Est...

-> Processing paper match [2606.02113]: A Primer in Post-Training Reasoning Data: What We ...

-> Processing paper match [2606.02111]: Jailbreaking Multimodal Large Language Models usin...

-> Processing paper match [2606.02109]: BADGER: Bridging Agentic and Deterministic Evaluat...

-> Processing paper match [2606.02107]: Network Distributed Multi-Agent Reinforcement Lear...

-> Processing paper match [2606.02106]: When Tabular Foundation Models Transfer Across Mod...

-> Processing paper match [2606.02101]: It does what it says on the tin: safe synthetic da...

-> Processing paper match [2606.02100]: PortBERT: Navigating the Depths of Portuguese Lang...

-> Processing paper match [2606.02093]: The Role of Ambiguity in Error Prediction via Unce...

-> Processing paper match [2606.02092]: LALE: Lightweight-Transformer Architecture for Lan...

-> Processing paper match [2606.02091]: DFlare: Scaling Up Draft Capacity for Block Diffus...

-> Processing paper match [2606.02080]: Agentic-J: An AI Agent for Biological Microscopy I...

-> Processing paper match [2606.02078]: Beyond ℓ_2 -norm and ℓ_∞ -norm: A Cur...

-> Processing paper match [2606.02073]: Planar Symmetric Pattern Generation...

-> Processing paper match [2606.02068]: Fast and Lightweight Novel View Synthesis with Dif...

-> Processing paper match [2606.02061]: Ablating Archetypes: The Stability of Archetypal S...

-> Processing paper match [2606.02060]: Where Do Deep-Research Agents Go Wrong? Span-Level...

-> Processing paper match [2606.02055]: Query-Limited Community Recovery in Stochastic Blo...

-> Processing paper match [2606.02054]: eMoT: evolving Memory-of-Thought via Symbolic Anch...

-> Processing paper match [2606.02049]: Explainable Data-driven Deep Reinforcement Learnin...

-> Processing paper match [2606.02048]: Topological texture analysis of microscopy images ...

-> Processing paper match [2606.02047]: Convex Distance Operator Transport: A Convex and G...

-> Processing paper match [2606.02045]: Attention mechanisms and transfer learning for rob...

-> Processing paper match [2606.02044]: Realistic noise synthesis reduces bias and improve...

-> Processing paper match [2606.02041]: SentGuard: Sentence-Level Streaming Guardrails for...

-> Processing paper match [2606.02038]: Uncertainty-Aware Graph Neural Reconstruction of U...

-> Processing paper match [2606.02035]: RL-ACRGNet: Reinforcement Learning-Based Chest Rad...

-> Processing paper match [2606.02031]: OpenWebRL: Demystifying Online Multi-turn Reinforc...

-> Processing paper match [2606.02027]: World-Task Factorization for Robot Learning...

-> Processing paper match [2606.02022]: Ranking vs. Assignment: The Metric Mismatch in Mul...

-> Processing paper match [2606.02020]: Unveiling the Entropy Dynamics of Chain-of-Thought...

-> Processing paper match [2606.02016]: Evaluating Real-World Generalizability of Algorith...

-> Processing paper match [2606.02011]: Extreme Low-Bit Inference in Reasoning Models: Fai...

-> Processing paper match [2606.02010]: PlanarBench: Evaluating LLM Spatial Reasoning via ...

-> Processing paper match [2606.02009]: Automated Essay Scoring and Language Certification...

-> Processing paper match [2606.02008]: Provable Data Scaling Law for Meta Learning via Co...

-> Processing paper match [2606.02004]: Machine Learning for Coding Retail Product Names t...

-> Processing paper match [2606.02001]: Scaling Agentic Capabilities via Grounded Interact...

-> Processing paper match [2606.02000]: Towards 3D-Aware Video Diffusion Models: Render-Fr...

-> Processing paper match [2606.01999]: Why Do Time Series Models Need Long Context Window...

-> Processing paper match [2606.01995]: CARTE: A Benchmark for Mapping Language Model Know...

-> Processing paper match [2606.01993]: MMG2Skill: Can Agents Distill In-the-Wild Guides i...

-> Processing paper match [2606.01992]: A Structured Benchmark for Text-Guided Anomaly Det...

-> Processing paper match [2606.01991]: SafeMCP: Proactive Power Regulation for LLM Agent ...

-> Processing paper match [2606.01987]: Graph Edit Distance Formulation for the Vehicle Ro...

-> Processing paper match [2606.01982]: An NLP-Driven Framework for Curriculum-Labor Marke...

-> Processing paper match [2606.01975]: Algorithmic algorithm development with LLMs: A Cas...

-> Processing paper match [2606.01973]: A Closer Look at In-Distribution vs. Out-of-Distri...

-> Processing paper match [2606.01967]: Training Prompt Matters: State-Adaptive Optimizati...

-> Processing paper match [2606.01964]: Eyettention II: A Dual-Sequence Architecture for M...

-> Processing paper match [2606.01961]: AutoMedBench: Towards Medical AutoResearch with Ag...

-> Processing paper match [2606.01954]: Flow-Transformed Implicit Processes for Function-S...

-> Processing paper match [2606.01952]: Randomized Least Squares Value Iteration itself is...

-> Processing paper match [2606.01950]: Learning Action-Conditional and Object-Centric Gau...

-> Processing paper match [2606.01948]: Rank-Constrained Deep Matrix Completion for Group ...

-> Processing paper match [2606.01947]: Parameter-Efficient Fine-Tuning of Large Pretraine...

-> Processing paper match [2606.01936]: What to Format and How: A Benchmark and Workflow A...

-> Processing paper match [2606.01934]: HMPO: Hybrid Median-length Policy Optimization for...

-> Processing paper match [2606.01929]: VET: A Framework for Analyzing AI Discourse...

-> Processing paper match [2606.01926]: Mitigating Bias in Locally Constrained Decoding vi...

-> Processing paper match [2606.01923]: Resonant Context Anchoring: Decoupling Attention R...

-> Processing paper match [2606.01914]: Mechanistic Diagnostics of Spatial Lexical Bias in...

-> Processing paper match [2606.01912]: SMH-Bench: Benchmarking LLM Agents for Environment...

-> Processing paper match [2606.01909]: Echo: A Joint-Embedding Predictive Architecture fo...

-> Processing paper match [2606.01908]: Private and Stable Test-Time Adaptation with Diffe...

-> Processing paper match [2606.01906]: Bayesian Spectral Emotion Transition Discovery fro...

-> Processing paper match [2606.01904]: KliniskVestBERT: BERT Model Specialised to Norwegi...

-> Processing paper match [2606.01901]: The Image Reconstruction Game: Drawing Common Grou...

-> Processing paper match [2606.01899]: RA-LWLM: Retrieval-Augmented In-Context Localizati...

-> Processing paper match [2606.01897]: Community-Aware Assessment of Social Textual Engag...

-> Processing paper match [2606.01896]: Train, Test, Re-evaluate: Schedule-Sensitive Evalu...

-> Processing paper match [2606.01895]: Collaborative Space Object Detection with Multi-Sa...

-> Processing paper match [2606.01894]: Physically-Constrained Mamba-SDE for Remaining Use...

-> Processing paper match [2606.01891]: MidSurfNet: Learnable Face Pairing and Interferenc...

-> Processing paper match [2606.01890]: Segment-driven Structural Induction and Semantic A...

-> Processing paper match [2606.01886]: Absorbing Complexity: An Interaction-Native Knowle...

-> Processing paper match [2606.01884]: EVA-Net: Subject-Independent EEG Motor Decoding wi...

-> Processing paper match [2606.01883]: Beyond the Simplex: Balanced Prototype Geometry fo...

-> Processing paper match [2606.01879]: CultureForest: Understanding and Evaluating Cultur...

-> Processing paper match [2606.01873]: G2LoRA: Gradient Orthogonal Low-Rank Adaptation Fr...

-> Processing paper match [2606.01869]: WorldCoder-Bench: Benchmarking Physically Grounded...

-> Processing paper match [2606.01868]: Task-Induced Representational Invariances Depend o...

-> Processing paper match [2606.01863]: Continual Learning as a Multiphase Moving-Boundary...

-> Processing paper match [2606.01862]: RadioMaster: Multi-Agent System for Autonomous Rad...

-> Processing paper match [2606.01861]: A Theoretical Framework for Self-Play Theorem Prov...

-> Processing paper match [2606.01856]: Boosting Multimodal Federated Learning via Chained...

-> Processing paper match [2606.01850]: Does Compression Preserve Uncertainty? A Unified B...

-> Processing paper match [2606.01849]: ContinuousBench: Can Differentially Private Synthe...

-> Processing paper match [2606.01847]: The Lie We Tell: Correcting the Euclidean Fallacy ...

-> Processing paper match [2606.01846]: Mos-Gen: A Generative Molecular Framework for Mosq...

-> Processing paper match [2606.01845]: Unveiling the Limits of Large Language Models in I...

-> Processing paper match [2606.01843]: Suppressing Forgery-Specific Shortcuts for General...

-> Processing paper match [2606.01840]: Evaluation of Baseline Methods for IDD-based SSD E...

-> Processing paper match [2606.01839]: Observation, Not Prediction: Conversation-Level Di...

-> Processing paper match [2606.01838]: LayerRoute: Input-Conditioned Adaptive Layer Skipp...

-> Processing paper match [2606.01834]: Physics-Guided Attention in a Lightweight TCN for ...

-> Processing paper match [2606.01833]: Learning Implicit Bias in Generative Spaces for Ac...

-> Processing paper match [2606.01830]: CAPF: Guiding Search-Agent Rollouts with Credit-At...

-> Processing paper match [2606.01828]: Dynamic Trust-Aware Sparse Communication Topology ...

-> Processing paper match [2606.01827]: Adaptive Sharpness-Aware Minimization with a Polya...

-> Processing paper match [2606.01820]: TalkTag: Fine-Grained Morphosyntactic Error Annota...

-> Processing paper match [2606.01816]: Site4Drug: Predicting Drug-Binding Target Sites wi...

-> Processing paper match [2606.01815]: CRAB-Bench: Evaluating LLM Agents under Complex Ta...

-> Processing paper match [2606.01813]: Cost-Aware Diffusion Draft Trees for Speculative D...

-> Processing paper match [2606.01811]: "I've Seen How This Goes": Characterizing Diversit...

-> Processing paper match [2606.01810]: Token Predictors Are Not Planners: Building Physic...

-> Processing paper match [2606.01806]: ProbeScale: Probing Analysis to Optimize Neural Sc...

-> Processing paper match [2606.01803]: OctoT2I: A Self-Evolving Agentic Text-to-Image Rou...

-> Processing paper match [2606.01802]: MOSS-Audio Technical Report...

-> Processing paper match [2606.01800]: Multilinguality of Large Language Models From a St...

-> Processing paper match [2606.01799]: Tree-Guided Identify-Then-Exploit: A Unified Frame...

-> Processing paper match [2606.01790]: STaR-KV: Spatio-Temporal Adaptive Re-weighting for...

-> Processing paper match [2606.01789]: Consistency evaluation of benchmarks used for caus...

-> Processing paper match [2606.01787]: Stochastic convergence of parallel asynchronous ad...

-> Processing paper match [2606.01783]: Breaking the Information Silo: Semantic Personas f...

-> Processing paper match [2606.01781]: Structure-Guided Adaptive Propagation for Protein-...

-> Processing paper match [2606.01779]: HarnessForge: Joint Harness and Policy Evolution f...

-> Processing paper match [2606.01774]: FLARE: Diffusion for Hybrid Language Model...

-> Processing paper match [2606.01770]: Adaptive Auto-Harness: Sustained Self-Improvement ...

-> Processing paper match [2606.01767]: EvoBrain: Continual Learning of EEG Foundation Mod...

-> Processing paper match [2606.01765]: An Algebraic View of the Expressivity of Recurrent...

-> Processing paper match [2606.01764]: Accelerating Min-Max Optimization via Power-Law St...

-> Processing paper match [2606.01755]: TriAlign: Towards Universal Truth Consistency in P...

-> Processing paper match [2606.01747]: Construction of Historical Knowledge Graphs Based ...

-> Processing paper match [2606.01746]: Sensitivity as a Double-Edged Sword: A Trade-off B...

-> Processing paper match [2606.01741]: SECUREVENT: Hybrid AI/ML Security Monitoring for D...

-> Processing paper match [2606.01738]: THRD: A Training-Free Multi-Turn Defense Framework...

-> Processing paper match [2606.01737]: TrafficRAG: A Multimodal RAG Framework for Traffic...

-> Processing paper match [2606.01736]: Argument Collapse: LLMs Flatten Long-Form Public D...

-> Processing paper match [2606.01734]: FlatVPR: Plug-and-play Geo-linear Residual Adapter...

-> Processing paper match [2606.01730]: Evidence-Gated LLM Priors for Multi-Objective Baye...

-> Processing paper match [2606.01725]: Characterization of Multi-Model Agentic AI Systems...

-> Processing paper match [2606.01723]: Shortcut to Nowhere: Demystifying Deep Spurious Re...

-> Processing paper match [2606.01722]: Post-Deterministic Distributed Systems: A New Foun...

-> Processing paper match [2606.01720]: A Note on Stability for Orthogonalized Matrix Mome...

-> Processing paper match [2606.01719]: Fair Finetuning Mitigates Distribution Inference A...

-> Processing paper match [2606.01717]: Decentralized Instruction Tuning: Conflict-Aware S...

-> Processing paper match [2606.01710]: Density-Aware Translation of Spurious Correlations...

-> Processing paper match [2606.01708]: Two-Fidelity Best-Action Identification for Stocha...

-> Processing paper match [2606.01703]: JenBridge: Adaptive Long-Form Video Soundtracking ...

-> Processing paper match [2606.01702]: KDH-CAD: Knowledge-data hybrid CAD learning under ...

-> Processing paper match [2606.01697]: RCEM: Embedder Equipped with Query Rewriting Skill...

-> Processing paper match [2606.01695]: CANARY: Zero-Label Detection of Fine-Tuning Contam...

-> Processing paper match [2606.01694]: Understanding Identity Continuity in Thermal Video...

-> Processing paper match [2606.01691]: IstGPT: LLM-based Anomaly Detection for Spatial-Te...

-> Processing paper match [2606.01689]: RPCASSM: Robust PCA State Space Model For Infrared...

-> Processing paper match [2606.01686]: HAIM: Human-AI Music Datasets for AI Music Product...

-> Processing paper match [2606.01682]: Off-the-Shelf LLMs as Process Scorers: Training-Fr...

-> Processing paper match [2606.01680]: Don't Let a Few Network Failures Slow the Entire A...

-> Processing paper match [2606.01679]: Encoded but Not Routed: Explaining the Table-Chart...

-> Processing paper match [2606.01678]: Why Do Self-Harm Prediction Models Struggle to Gen...

-> Processing paper match [2606.01672]: RDA: Reward Design Agent for Reinforcement Learnin...

-> Processing paper match [2606.01671]: When Meaning Travels: A Granular Lens on Hybrid-Mo...

-> Processing paper match [2606.01670]: Time-Aware Diffusion based on Preference Disentang...

-> Processing paper match [2606.01667]: ATLAS: Agentic Test-time Learning-to-Allocate Scal...

-> Processing paper match [2606.01666]: DOT-MoE: Differentiable Optimal Transport for MoEf...

-> Processing paper match [2606.01665]: Quantifying the Energy Floor: Direct Measurement a...

-> Processing paper match [2606.01660]: Gate the Filter, Not the Message: Node-Channel Mix...

-> Processing paper match [2606.01655]: MINTS: Minimalist Thompson Sampling...

-> Processing paper match [2606.01645]: Self-Regulating Annealing in Heavy-Tailed Diffusio...

-> Processing paper match [2606.01640]: MobEvolve: An Agentic Self-Evolving Heuristic Syst...

-> Processing paper match [2606.01637]: Easier to Mislead Than to Correct: Harmful and Ben...

-> Processing paper match [2606.01635]: AlphaToken: Decoupling Adaptation and Stability fo...

- > Processing paper match [2606.01634]: E4GEN: Event-level Explainable Extreme-Enhanced Ti...
- > Processing paper match [2606.01632]: A Framework for Graph-Conditioned Hierarchical Sha...
- > Processing paper match [2606.01629]: Benchmarking LLM-as-a-Judge for Long-Form Output E...
- > Processing paper match [2606.01628]: Demystifying Multimodal Biomolecular Co-design Wit...
- > Processing paper match [2606.01626]: IMWM: Intuition Models Complement World Models for...
- > Processing paper match [2606.01619]: ReSkill: Reconciling Skill Creation with Policy Op...
- > Processing paper match [2606.01617]: EvoPool: Evolutionary Programmatic Annotation for ...
- > Processing paper match [2606.01613]: TechGraphRAG: An Agentic Graph-Augmented RAG Frame...
- > Processing paper match [2606.01612]: Self-Improving Small Object Grounding in LVLMS...
- > Processing paper match [2606.01610]: Revisiting Ripple Effects in Knowledge Editing thr...
- > Processing paper match [2606.01607]: FedMTFI: Feature Importance Based Optimized Multi ...
- > Processing paper match [2606.01602]: Estimating Mutual Information between Time Series ...
- > Processing paper match [2606.01600]: RoboTrustBench: Benchmarking the Trustworthiness o...
- > Processing paper match [2606.01599]: TRON: Targeted Rule-Verifiable Online Environments...
- > Processing paper match [2606.01596]: Learning Chaotic Dynamics through Second-Order Geo...
- > Processing paper match [2606.01595]: Uncertainty-Calibrated Diffusion for Reliable 3D M...
- > Processing paper match [2606.01591]: TLG: Temporal-Logic Grounding for Video Question A...
- > Processing paper match [2606.01584]: Identifying High-Confidence Social Biases in LLMs ...
- > Processing paper match [2606.01567]: Defenses & Enablers For Skill Injection Attacks on...
- > Processing paper match [2606.01566]: RobustModelMaker: Coupling Bootstrap Stability Sel...
- > Processing paper match [2606.01563]: MomentKV: Closing the Directional Gap in KV Cache ...
- > Processing paper match [2606.01561]: S-SPPO: Semantic-Calibrated Self-Play Preference O...
- > Processing paper match [2606.01560]: GJDNet: Robust Graph Neural Networks via Joint Dis...
- > Processing paper match [2606.01557]: Everywhere Learning: Artificial Intelligence with ...
- > Processing paper match [2606.01552]: RoleCDE: Benchmarking and Mitigating Role-Alignment...
- > Processing paper match [2606.01546]: Flexible Online Representation Learning Based on S...

- > Processing paper match [2606.01544]: CRePE: Convolution-aware Relative Importance in Po...
- > Processing paper match [2606.01542]: Self-Conditioned Positional HNSW for Overlap-Aware...
- > Processing paper match [2606.01540]: TN-SHAP-G: Graph-Structured Tensor Network Surroga...
- > Processing paper match [2606.01539]: Scalable Counterfactual Risk Estimation for Rare E...
- > Processing paper match [2606.01538]: MPMWorlds: Material-Point-Method Simulations for I...
- > Processing paper match [2606.01537]: PaCX-MAE: Physiology-Augmented Chest X-Ray Masked ...
- > Processing paper match [2606.01533]: Multi-Agent Computer Use...
- > Processing paper match [2606.01532]: Rethinking the Role of Positional Encoding: Slidin...
- > Processing paper match [2606.01528]: Joint Agent Memory and Exploration Learning via No...
- > Processing paper match [2606.01527]: Near-Optimal Pure Machine Unlearning for Smooth St...
- > Processing paper match [2606.01525]: Semi-Supervised Hyperbolic Hierarchical Clustering...
- > Processing paper match [2606.01521]: Fast Generalization after Interpolation via Critic...
- > Processing paper match [2606.01520]: TERRA: Task-Embedded Reasoning and Representation ...
- > Processing paper match [2606.01513]: Compliance-Scored Best-of-N Guardrail Orchestratio...
- > Processing paper match [2606.01509]: ProbMoE: Differentiable Probabilistic Routing for ...
- > Processing paper match [2606.01508]: Agent Operating Systems (AOS): Integrating Agentic...
- > Processing paper match [2606.01504]: Semantic Retrieval for Product Search in E-Commerc...
- > Processing paper match [2606.01503]: On the Limits of Token Reduction for Efficient Uni...
- > Processing paper match [2606.01502]: Move the Query, Not the Cache: Characterizing Cros...
- > Processing paper match [2606.01498]: TimeSage-MT: A Multi-Turn Benchmark for Evaluating...
- > Processing paper match [2606.01495]: CART: Context-Anchored Recurrent Transformer -- A ...
- > Processing paper match [2606.01494]: ClawHub Security Signals: When VirusTotal, Static ...
- > Processing paper match [2606.01490]: LLM Consortium for Software Design Refinement: A C...
- > Processing paper match [2606.01485]: Perception First: A Frontier Native-Video Model wi...
- > Processing paper match [2606.01483]: MURMUR: An Efficient Inference System for Long-For...
- > Processing paper match [2606.01482]: Beyond Topical Similarity: Contrastive Evidence Re...

-> Processing paper match [2606.01479]: Sparse Autoencoders for Interpretable Emotion Cont...

-> Processing paper match [2606.01478]: Crazyflow: An Accurate, GPU-Accelerated, Different...

-> Processing paper match [2606.01476]: OmniOPD: Logit-Free On-Policy Distillation via Spe...

-> Processing paper match [2606.01473]: A Minimalist Brain-Computer Musical Interface for ...

-> Processing paper match [2606.01472]: Hierarchical Online Prompt Mutation with Dual-Loop...

-> Processing paper match [2606.01470]: Emergent Transfer of a Physics Foundation Model fr...

-> Processing paper match [2606.01469]: Peacemaker at ATE-IT: Automatic term extraction fr...

-> Processing paper match [2606.01468]: Computation-Aware Kalman Filtering with Model Sele...

-> Processing paper match [2606.01464]: Cross-lingual Self-Consistency for Multilingual Re...

-> Processing paper match [2606.01462]: An Enigma of Artificial Reason: Investigating the ...

-> Processing paper match [2606.01461]: Genotype-Conditioned Molecular Generation via Evid...

-> Processing paper match [2606.01457]: Transferring Information Across Interventions in C...

-> Processing paper match [2606.01456]: Truthful AI Advisors: A Pre-Specified Benchmark fo...

-> Processing paper match [2606.01451]: Before and After Temperature: A Distributional Vie...

-> Processing paper match [2606.01446]: Spatially Distributed Task-Oriented Compression fo...

-> Processing paper match [2606.01444]: Self-Revising Discovery Systems for Science: A Cat...

-> Processing paper match [2606.01443]: UR-JEPA: Uniform Rectifiability as a Regularizer f...

-> Processing paper match [2606.01442]: On the Evaluation of Spiking Neural Network Config...

-> Processing paper match [2606.01441]: Dive into Ambiguity: A*-Inspired Multi-Agents Comm...

-> Processing paper match [2606.01437]: CEAR: Certified Ensemble Adversarial Robustness in...

-> Processing paper match [2606.01436]: Learning from Saturated Data: Signals Beyond Corre...

-> Processing paper match [2606.01435]: Don't Ask the LLM to Track Freshness: A Determinis...

-> Processing paper match [2606.01434]: DrugClaw and DrugAudit: A Primary-Source-Grounded ...

-> Processing paper match [2606.01432]: Leaf Spectral Reflectance Prediction Using Multi-H...

-> Processing paper match [2606.01427]: On the Uncertainty Quantification Ability of Tabul...

-> Processing paper match [2606.01425]: Learning-based Directed Graph Abstraction of Combi...

-> Processing paper match [2606.01421]: Target localization, identification and sensing us...

-> Processing paper match [2606.01417]: GovAI-Pipe: A Layered AI Governance Pipeline for C...

-> Processing paper match [2606.01416]: Self-Healing Agentic Orchestrators for Reliable To...

-> Processing paper match [2606.01413]: Differentially Private Datastore Generation for Re...

-> Processing paper match [2606.01412]: GPTQ-intrinsic LoRA: A Near-optimal Algorithm for ...

-> Processing paper match [2606.01402]: Neural Network Compression by Approximate Differen...

-> Processing paper match [2606.01400]: Consistent and Distinctive: LLM Benchmark Efficien...

-> Processing paper match [2606.01397]: Autopilot-Preserving Residual Q-Learning with HJB-...

-> Processing paper match [2606.01394]: UniD³: A Knowledge Graph-Enhanced RAG Framework...

-> Processing paper match [2606.01393]: Dr. DocBench: A Comprehensive Benchmark for Expert...

-> Processing paper match [2606.01386]: GuidaPA: Privacy-Preserving Chatbot for Public Adm...

-> Processing paper match [2606.01385]: Bridging Requirements and Architecture: Multi-Agen...

-> Processing paper match [2606.01382]: Efficient Exploration for Iterative Nash Preferenc...

-> Processing paper match [2606.01379]: Turning Back Without Forgetting: Selective Backwar...

-> Processing paper match [2606.01375]: Beyond Access: Guided LLM Scaffolding for Independ...

-> Processing paper match [2606.01374]: From Performance to Viability: A Bootstrap Framewo...

-> Processing paper match [2606.01372]: BRo-JEPA: Learning Modular Arithmetic in Latent Sp...

-> Processing paper match [2606.01365]: Early Diagnosis of Wasted Computation in Multi-Age...

-> Processing paper match [2606.01364]: Needles at Scale: LLM-Assisted Target Selection fo...

-> Processing paper match [2606.01363]: All Models are Wrong, Knowing Where is Useful: On ...

-> Processing paper match [2606.01352]: FlowTime: Towards Continuous Generative Watch Time...

-> Processing paper match [2606.01351]: Recognize Your Orchestrator: An Entropy Dynamics P...

-> Processing paper match [2606.01342]: Towards Optimal Robustness in Learning-Augmented P...

-> Processing paper match [2606.01340]: Sample Complexity and Decision-Theoretic Guarantee...

-> Processing paper match [2606.01339]: FreqLite: A Lightweight Frequency-Decomposed Linea...

-> Processing paper match [2606.01338]: Benchmarking Local LLMs for Natural-Language-to-SQ...

-> Processing paper match [2606.01336]: LongAttnComp: Cross-Family Context Compression for...

-> Processing paper match [2606.01329]: Conditioned free-energy density of proteins using ...

-> Processing paper match [2606.01325]: SEArch: Optimistic Policy Selection Between Scene ...

-> Processing paper match [2606.01324]: Digital Twin-Assisted Adaptive Multi-Agent DRL for...

-> Processing paper match [2606.01323]: DiffuSent: Towards a Unified Diffusion Framework f...

-> Processing paper match [2606.01322]: TukaBench: A Culturally Grounded Jailbreak Benchma...

-> Processing paper match [2606.01316]: Science Earth: Towards A Planet-Scale Operating Sy...

-> Processing paper match [2606.01314]: SkillSmith: Co-Evolving Skills and Tools for Self...

-> Processing paper match [2606.01313]: PSG-Nav: Probabilistic Scene Graph Navigation via ...

-> Processing paper match [2606.01312]: A Communication-Centric 6G-LLM Architecture for Sc...

-> Processing paper match [2606.01311]: SkillAdaptor: Self-Adapting Skills for LLM Agents ...

-> Processing paper match [2606.01306]: FAiT: Frequency-Aware Inverted Transformer for Mul...

-> Processing paper match [2606.01304]: When Hard Negatives Hurt: Bridging the Generative-...

-> Processing paper match [2606.01302]: Structure and Scale in Simplicial Sequence Modelli...

-> Processing paper match [2606.01301]: Med-HEAL: Analyzing and Mitigating Hallucinations ...

-> Processing paper match [2606.01300]: ChronosAD: Leveraging Time Series Foundation Model...

-> Processing paper match [2606.01298]: Challenger at MultiPRIDE: Is It Hate Speech or Rec...

-> Processing paper match [2606.01294]: Don't Read Everything: A Curvature-Conditioned Que...

-> Processing paper match [2606.01293]: ResNet-34 with Lightweight Decoder for Accurate an...

-> Processing paper match [2606.01292]: What Makes a Strong Model? A Unified Spectral Anal...

-> Processing paper match [2606.01291]: Quantum Algorithm for Distributed Reduction of Ent...

-> Processing paper match [2606.01289]: Feature to Dynamics: Feature-space to Autoregressi...

-> Processing paper match [2606.01287]: Beyond Visual Memory: Mechanistic Diagnostics of L...

-> Processing paper match [2606.01286]: BenchEvolver: Frontier Task Synthesis via Solution...

-> Processing paper match [2606.01285]: Knowledge-Intensive Video Generation...

-> Processing paper match [2606.01283]: AdaKernel: Learning Adaptive Kernel Parameters for...

- > Processing paper match [2606.01282]: KG-FairDiff: Knowledge Graph-Guided Prompt Refinem...
- > Processing paper match [2606.01281]: RLVR without Ineffective Samples: Group Prioritize...
- > Processing paper match [2606.01279]: ANDES: Agent Native Data Evolving Synthesis Tool f...
- > Processing paper match [2606.01277]: DeepIPCv3: Event-Aware Multi-Modal Sensor Fusion f...
- > Processing paper match [2606.01276]: Worlds Within Words: Translating Culture in Ancien...
- > Processing paper match [2606.01273]: GLIDE: Graph-guided Leap Inference for Diffusion E...
- > Processing paper match [2606.01269]: Emergent Ordinal Geometry in Transformers Trained ...
- > Processing paper match [2606.01265]: PALTO: Physics-Informed Active Learning for Tri-Ga...
- > Processing paper match [2606.01260]: IndoBias: A Dual Track Culturally Grounded Benchma...
- > Processing paper match [2606.01258]: Beyond Sinusoids: A Morlet Wavelet Framework for T...
- > Processing paper match [2606.01256]: Distribution-free changepoint localization after s...
- > Processing paper match [2606.01255]: Agentic Clustering: Controllable Text Taxonomies v...
- > Processing paper match [2606.01252]: Understanding LLM Behavior in Multi-Target Cross-L...
- > Processing paper match [2606.01249]: Trust Region On-Policy Distillation...
- > Processing paper match [2606.01246]: SIRIUS-SQL: Anchoring Multi-Candidate Text-to-SQL ...
- > Processing paper match [2606.01244]: Efficient Approximation for Encoder--Decoder Neura...
- > Processing paper match [2606.01243]: Unlocking the Black Box of Latent Reasoning: An In...
- > Processing paper match [2606.01240]: Efficient RAG with Intent-Aware Retrieval and Sema...
- > Processing paper match [2606.01238]: Training-Free Imitation Learning with Closed-Form ...
- > Processing paper match [2606.01237]: Brain-Atlas-Guided Generative Counterfactual Atten...
- > Processing paper match [2606.01234]: Differing Roles of Leisure and Productivity in GDP...
- > Processing paper match [2606.01230]: HomeFlow: A Data Flywheel for Smart Home Agent Tra...
- > Processing paper match [2606.01229]: Application of Algorithms in Energy-Efficient Desi...
- > Processing paper match [2606.01227]: DAGGER: Gradient-Free Construction of Transiently ...
- > Processing paper match [2606.01224]: Advanced Mathematics Learning Behavior Prediction ...
- > Processing paper match [2606.01223]: Connecting the Dots: Benchmarking Reflective Memor...

- > Processing paper match [2606.01221]: Hybrid Imbalanced Regression Through Unified Data-...
- > Processing paper match [2606.01220]: Fine-Tuning Diffusion Models for Molecular Generat...
- > Processing paper match [2606.01217]: Analysis of Ethnic Disparities in Autism Spectrum ...
- > Processing paper match [2606.01216]: Riemannian Optimization for Hadamard Products of L...
- > Processing paper match [2606.01215]: Distilling Neuro-Symbolic Programs into 3D Multi-m...
- > Processing paper match [2606.01213]: TECCI: Tricky Edits of Collected and Curated Image...
- > Processing paper match [2606.01212]: DiscourseFlip: An Oblique Discourse-Level Opinion ...
- > Processing paper match [2606.01207]: Feature Alignment Determines Fusion Strategy: A Co...
- > Processing paper match [2606.01204]: Implicit Geographic Inference in LLM Medical Triag...
- > Processing paper match [2606.01202]: The Shape of Wisdom: Decision Trajectories in Lang...
- > Processing paper match [2606.01199]: Can LLM Agents Sustain Long-Horizon Organizational...
- > Processing paper match [2606.01198]: Linear Strategic Classification with Endogenous Im...
- > Processing paper match [2606.01196]: Low-Resource Safety Failures Are Action Failures, ...
- > Processing paper match [2606.01189]: The Case for Model Science: Verify, Explore, Steer...
- > Processing paper match [2606.01188]: pcbGPT: Automatic PCB Schematic Synthesis from Nat...
- > Processing paper match [2606.01185]: "Skill issues": data-centric optimization of lake...
- > Processing paper match [2606.01184]: Topological Ignorability for Structural Causal Eff...
- > Processing paper match [2606.01182]: CA-BED: Conversation-Aware Bayesian Experimental D...
- > Processing paper match [2606.01179]: Physics-Informed Deep Learning for Entropy Predict...
- > Processing paper match [2606.01176]: Temporal Motif Signatures for Temporal Graph Neura...
- > Processing paper match [2606.01172]: Revisiting Neural Processes via Fourier Transform ...
- > Processing paper match [2606.01171]: AI From the Margins (AIM): Rethinking Participator...
- > Processing paper match [2606.01168]: Thinking Economically: A Hierarchical Framework fo...
- > Processing paper match [2606.01166]: BraveGuard: From Open-World Threats to Safer Compu...
- > Processing paper match [2606.01162]: Deft Scheduling of Dynamic Cloud Workflows with Va...
- > Processing paper match [2606.01160]: Expected Value Alignment for Generative Reward Mod...

-> Processing paper match [2606.01159]: Fairness in two-player zero-sum games with bandit ...
-> Processing paper match [2606.01155]: When Data Is Scarce: Scaling Sparse Language Model...
-> Processing paper match [2606.01152]: ASE-26: a curriculum for agentic software engineer...
-> Processing paper match [2606.01151]: Lagrangian Perturbation Diffusion Steering: Latent...
-> Processing paper match [2606.01148]: Not All Explanations Simulate Equally: Comparing V...
-> Processing paper match [2606.01145]: Reasoning4Sciences: Bridging Reasoning Language Mo...
-> Processing paper match [2606.01139]: SkillRevise: Improving LLM-Authored Agent Skills v...
-> Processing paper match [2606.01138]: AMP: A Vendor-Neutral Wire Format for Agent Memory...
Successfully processed and generated 2743 text database chunks.

>>> Running Evaluation Matrix for Configuration: [FULL_AGENT] <<<

-> Solving Question ID: q01
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q02
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q03
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q04
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q05
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q06
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q07
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q08
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q09
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q10
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q11
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q12
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q13
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q14
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q15
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q16
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q17

Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q18
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q19
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q20
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q21
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q22
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q23
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q24
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q25
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q26
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q27
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q28
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q29
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q30
Pacing API requests... resting for 5 seconds...
Finished! Prediction file safely saved to: predictions/full_agent.jsonl

>>> Running Evaluation Matrix for Configuration: [BASELINE] <<<
-> Solving Question ID: q01
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q02
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q03
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q04
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q05
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q06
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q07
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q08
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q09
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q10
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q11
Pacing API requests... resting for 5 seconds...

-> Solving Question ID: q12
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q13
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q14
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q15
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q16
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q17
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q18
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q19
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q20
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q21
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q22
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q23
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q24
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q25
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q26
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q27
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q28
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q29
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q30
Pacing API requests... resting for 5 seconds...
Finished! Prediction file safely saved to: predictions/baseline.jsonl

>>> Running Evaluation Matrix for Configuration: [NO_PLANNER] <<<

-> Solving Question ID: q01
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q02
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q03
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q04
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q05
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q06

Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q07
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q08
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q09
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q10
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q11
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q12
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q13
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q14
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q15
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q16
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q17
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q18
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q19
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q20
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q21
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q22
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q23
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q24
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q25
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q26
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q27
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q28
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q29
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q30
Pacing API requests... resting for 5 seconds...
Finished! Prediction file safely saved to: predictions/no_planner.jsonl

>>> Running Evaluation Matrix for Configuration: [NO_REFLECTOR] <<<

[illegible]

-> Solving Question ID: q27
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q28
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q29
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q30
Pacing API requests... resting for 5 seconds...
Finished! Prediction file safely saved to: predictions/no_reflector.jsonl

>>> Running Evaluation Matrix for Configuration: [NO_HYBRID] <<<

-> Solving Question ID: q01
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q02
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q03
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q04
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q05
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q06
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q07
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q08
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q09
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q10
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q11
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q12
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q13
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q14
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q15
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q16
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q17
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q18
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q19
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q20
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q21

Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q22
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q23
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q24
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q25
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q26
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q27
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q28
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q29
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q30
Pacing API requests... resting for 5 seconds...
Finished! Prediction file safely saved to: predictions/no_hybrid.jsonl

>>> Running Evaluation Matrix for Configuration: [NO_VERIFIER] <<<

-> Solving Question ID: q01
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q02
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q03
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q04
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q05
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q06
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q07
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q08
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q09
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q10
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q11
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q12
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q13
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q14
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q15
Pacing API requests... resting for 5 seconds...

-> Solving Question ID: q16
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q17
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q18
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q19
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q20
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q21
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q22
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q23
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q24
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q25
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q26
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q27
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q28
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q29
Pacing API requests... resting for 5 seconds...
-> Solving Question ID: q30
Pacing API requests... resting for 5 seconds...
Finished! Prediction file safely saved to: predictions/no_verifier.jsonl

=====
ALL ABLATION FILES SUCCESSFULLY COMPILED FOR GRADING!
Check the 'predictions/' folder to review your outputs.
=====

(myenv) asifkhan@ASIFs-MacBook-Air Agentic_AI_Deep_Research %

(myenv) asifkhan@ASIFs-MacBook-Air Agentic_AI_Deep_Research % python3 src/
evaluator.py

=====
ADRA SYSTEM ABLATION EVALUATION MATRIX
=====

[FULL_AGENT HIGHLIGHTS]:

- > Observable Status: Complete (30/30 rows)
- > Calculated Accuracy: 92.2%
- > Empirical Citation Density: 2.6 papers/Q
- > Generated Character Volume: 1156 chars

[BASELINE HIGHLIGHTS]:

- > Observable Status: Complete (30/30 rows)
- > Calculated Accuracy: 88.5%
- > Empirical Citation Density: 1.2 papers/Q
- > Generated Character Volume: 930 chars

[NO_PLANNER HIGHLIGHTS]:

- > Observable Status: Complete (30/30 rows)
- > Calculated Accuracy: 91.3%
- > Empirical Citation Density: 1.6 papers/Q
- > Generated Character Volume: 984 chars

[NO_REFLECTOR HIGHLIGHTS]:

- > Observable Status: Complete (30/30 rows)
- > Calculated Accuracy: 91.5%
- > Empirical Citation Density: 1.7 papers/Q
- > Generated Character Volume: 833 chars

[NO_HYBRID HIGHLIGHTS]:

- > Observable Status: Complete (30/30 rows)
- > Calculated Accuracy: 92.8%
- > Empirical Citation Density: 2.5 papers/Q
- > Generated Character Volume: 924 chars

[NO_VERIFIER HIGHLIGHTS]:

- > Observable Status: Complete (30/30 rows)
- > Calculated Accuracy: 92.0%
- > Empirical Citation Density: 2.8 papers/Q
- > Generated Character Volume: 1174 chars

(myenv) asifkhan@ASIFs-MacBook-Air Agentic_AI_Deep_Research %